

**BEGINNING OF CLASS:**  
**GRAB A PIECE OF CARDSTOCK AND CREATE A NAME TENT**  
**WITH THE FIRST NAME YOU'D LIKE TO GO BY - write BIG**



Stand name tents on each desk

Join our iClicker class:

<https://join.iclicker.com/XSFZ>

## LECTURE 1

# Intro to CSCI 3022

Intro to the data science lifecycle and exploratory data analysis with Pandas

**CSCI 3022 @ CU Boulder**

Maribeth Oscamou

Content credit: [Acknowledgments](#)



## Lesson Learning Objectives:

- **Meet your classmates**
- **Name and explain the stages of the data science lifecycle**
- **Identify 5 key data properties to consider when doing Exploratory Data Analysis and implement using demo data**
- **Read in data from a .csv file to a Pandas DataFrame**

# Roadmap

---

Lecture 01, CSCI 3022

- [Intros & Logistics](#)
- [What is Data Science?](#)
- [Exploratory Data Analysis & Wrangling](#)
- Intro to Pandas
  - [Jupyter Notebook Demo: EDA using Pandas](#)
- [Supporting Materials:](#)
  - Intro to Pandas Data Structures

In Groups of 4 INTRODUCE YOURSELF :

- NAME, YEAR, MAJOR, HOMETOWN
- HOBBIES/INTERESTS
- SOME RANDOM FUN FACT ABOUT YOU

# Getting To Know You:

I'd like to get a chance to be introduced to each of you!

1. **Please sign-up for a 15 min. timeslot here: [link](#)**  
(this link is also in the Week 1 Info announcements on Canvas and Piazza)  
to meet with me during the first couple weeks to briefly introduce yourself and meet a few other classmates.

# Meet The Course Team



Isabella Longo  
Course Manager



Vincent Bowen  
Course Manager



Kevin Buhler  
Course Assistant



Grace Mudd  
Course Assistant



Noah Turner  
Course Assistant



Owen Vangermeersch  
Course Assistant

## Course Online platforms

---

**Canvas** (<https://canvas.colorado.edu/courses/117881> )

- Where all course information including lectures, assignments, announcements and grades are posted

**CSCI 3022 JupyterHub** (<https://coding.csel.io/> )

- Where you will work on all assignments (links on Canvas assignments automatically take you here).

**Piazza** (linked from Canvas)

- How to contact professor and discuss questions with other students

**Gradescope** (linked from Canvas)

- Where all assignments are submitted

**iClicker** (<https://student.iclicker.com/#/login>)

- Where you answer polls during class

# Accessing Lecture Slides and Jupyter Notebook

Canvas - Modules - Lesson Materials

<https://canvas.colorado.edu/courses/117881/modules>

2025 Spring Term

[Home](#)

[Announcements](#)

**[Modules](#)**

[Assignments](#)

[Gradescope](#)


[Piazza](#)

[Grades](#)


[NameCoach](#)

[Web Grading Sync](#)


[Quizzes](#) 


[Syllabus](#) 



[Discussions](#) 




[Pages](#) 


[Outcomes](#) 




 ▾ Lesson Materials - Week 1




 Lesson 1 - Intro to Data Science Lifecycle




  **[Lesson 1 - In Class Slides: Intro to Data Science Lifecycle](#)**

  [Lesson 1 - In Class Jupyter Demo](#) 

 Lesson 1 Supporting Materials

  [Reading - The Data Science Lifecycle \(LDS 1.1-1.3\)](#) 

  [Reading: Pandas DataFrames \(LDS: Chapter 6 Intro\)](#) 

  [Reading: Data Scope & Question \(LDS 6.1.1 & 6.1.2\)](#) 

# Course Logistics: Your *Typical* Week At A Glance

---

Mon	Tues	Wed	Thurs	Fri
Attend & Participate in Class		Attend & Participate in Class	HW Due 11:59pm via Gradescope	In Class Quiz (beginning of class) Attend & Participate in Class
Previous week quiz grades posted		Previous week HW grades posted		Next week HW released



# Course Logistics: Your First Week At A Glance

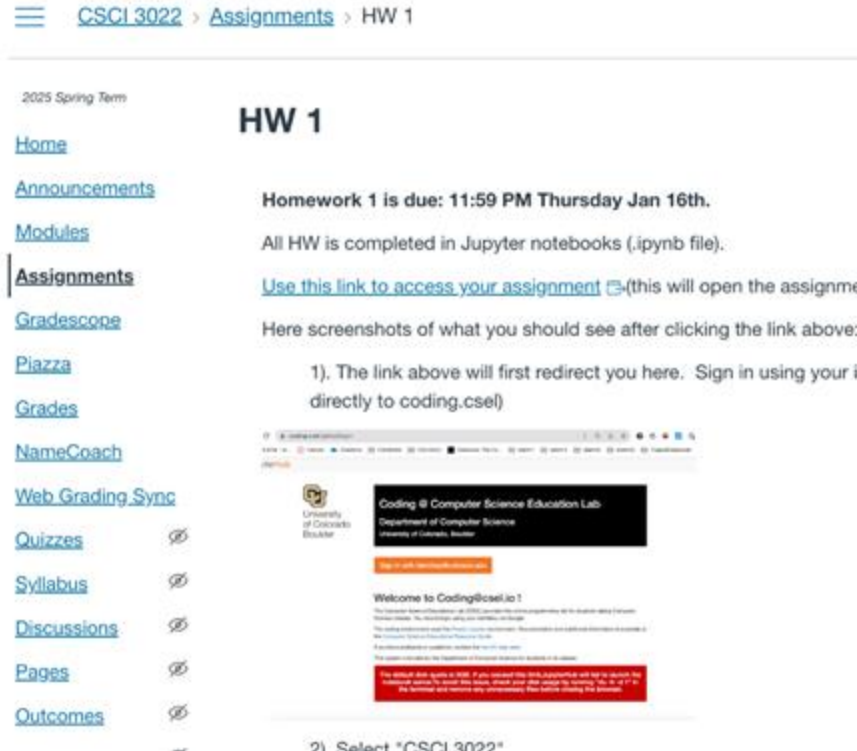
Mon 1/13	Tues 1/14	Wed 1/15	Thurs 1/16	Fri 1/17	
Attend & Participate in Class		Attend & Participate in Class		Attend & Participate in Class	
Office Hours Begin (See Schedule on Canvas)			HW 1 Due 11:59pm via Gradescope (Includes Intro to CSCI 3022 Video assignment)	<del>In-Class Quiz (beginning of class)</del>	
				HW 2 released	



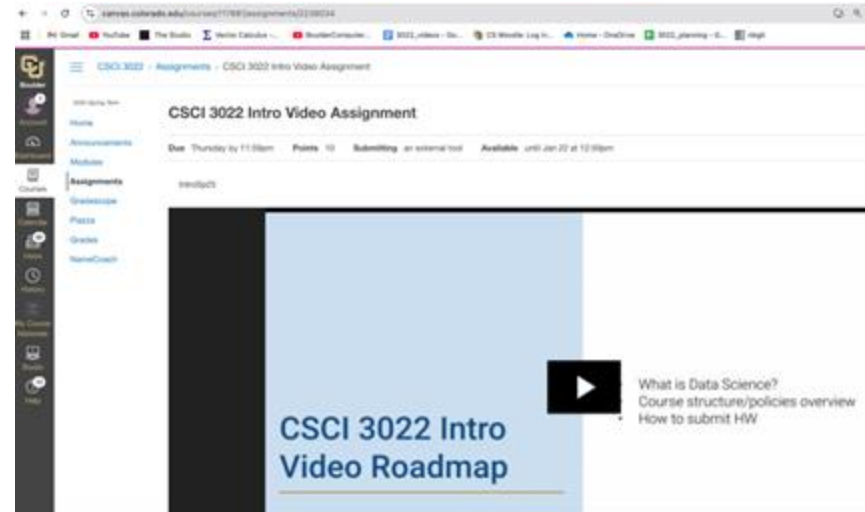
## Accessing HW 1

## Canvas -> Assignments

<https://canvas.colorado.edu/courses/117881/assignments/2238036>



## HW 1 Includes an Intro Video Assignment



# Office Hours:

---

<https://canvas.colorado.edu/courses/117881/pages/hw-slash-office-hours>

## Jupyter Notebook and LaTeX Troubleshooting and Tips

*Make sure before submitting to double check that your PDF includes all of the manually graded questions and plots, and that all code is fully visible in your PDF.*

### General best practices

- Make sure you have not renamed the .ipynb file. For example, HW 2 must be named hw02.ipynb
- Make sure you haven't inserted any new cells into the notebook.
- Make sure that you're in the 3022 instance of CSEL DataHub. You can do this by signing out of JupyterHub and then re-clicking the link. It should lead you to the page where you have to select the course "3022". The 3022 course has otter-grader installed in it. Other courses in the DataHub may not.
- If you make changes in your HW and run your export cell in your notebook more than once you should first delete the PDF (in the folder where the notebook is) and then re-run. It's possible that the version you submit is an earlier version of your HW.

### First fixes to try

- Save everything, delete the zip and pdf files and shut your browser window. Then open a new browser window and then restart your kernel and run through all of the cells and SAVE the nb before running the final export cell.
- As an extension, log out of coding.csel completely (after saving any work), close your browser, then launch a new one. Make sure you have selected CSCI 3022 as your coding environment.

### Latex Issues

- Check that there aren't any spaces after your dollar signs in LaTeX:

<https://docs.google.com/document/d/1ndr3Wj1PSF5qzILMaBJznwh6QGeEXjd5TAJ6nf9EJvo/edit?usp=sharing>

## Course Prerequisites (minimum grade C-)

---

- Data Structures (CSCI 2270 or equivalent)
- Calculus 2 (APPM 1360 or MATH 2300 or equivalent)
- Discrete Math (CSCI 2824 or equivalent)

# What is Data Science?


---

Lecture 01, CSCI 3022

- Intros & Logistics
- **What is Data Science?**
- Exploratory Data Analysis & Wrangling
- Pandas Data Structures
- Jupyter Notebook Demo: EDA using Pandas

# Some examples of Data Science?

# Recommendation Systems



# DJ Patil calls data scientists ‘a new kind of first responder’

By Rachel Leven | April 13, 2023

<https://www.youtube.com/watch?v=LiHMrn2AHpw>




DJ Patil spoke to UC Berkeley data science students on April 10. (Photo/

On March 14, 2020, the United States was on the brink of a pandemic. Covid-19 had killed at least 60 people and two cruise ships with ill passengers were set to dock in San Francisco. That's when DJ Patil received a call: How can data help California combat this?

So the former White House chief data scientist put together a plan. His team acquired hospital and community data, developed surveys, models, dashboards and data catalogs, and used those tools and insights to inform public officials across the state and the country.

"All of this came together in this effort to really take on Covid," said Patil, who is now a general partner at GreatPoint Ventures, at an April 10 UC Berkeley



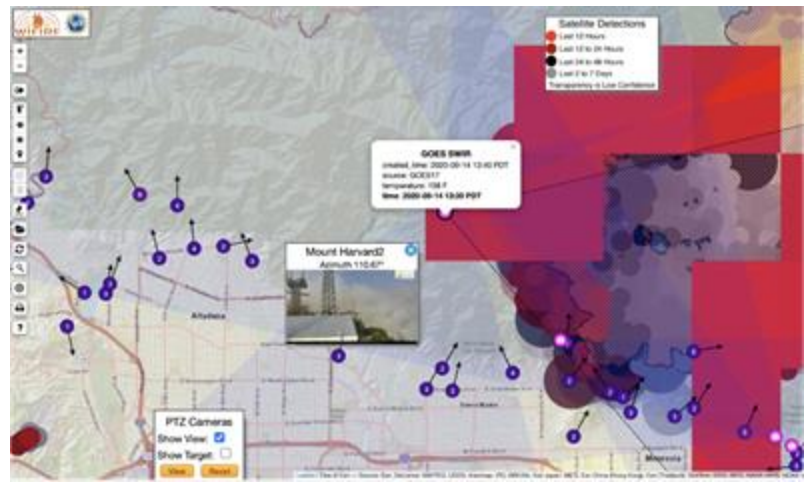
First U.S. Chief  
Data Scientist  
(Obama Adm.)

<https://www.sciencefriday.com/segments/an-exit-interview-with-u-s-chief-data-scientist-dj-patil/>



# WIFIRE (UCSD) - Wildfire modeling and management

<https://wifire.ucsd.edu>



# 2019: First Image of a Black Hole



Katie Bouman  
MIT/Caltech



Talk Video: <https://youtu.be/TSgpliktqwc>

## THE ASTROPHYSICAL JOURNAL LETTERS

### First M87 Event Horizon Telescope Results. III. Data Processing and Calibration

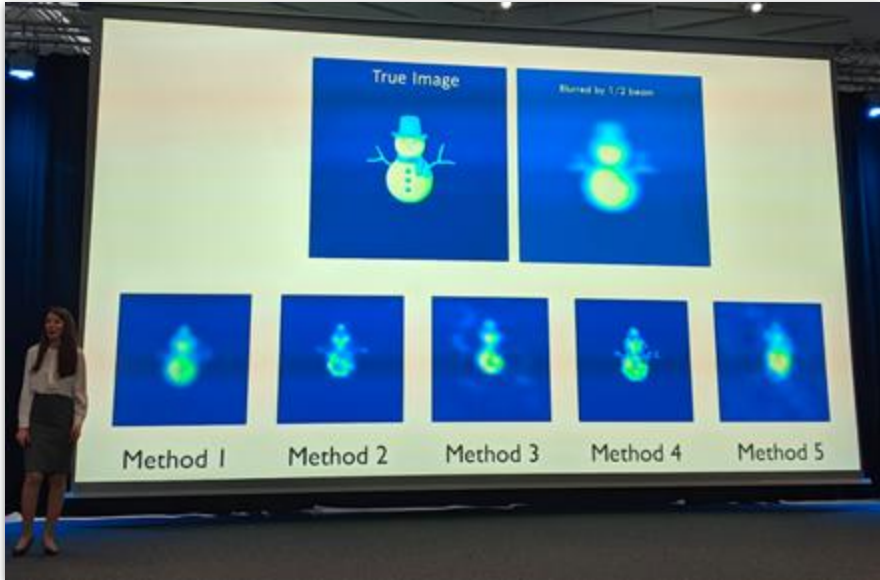
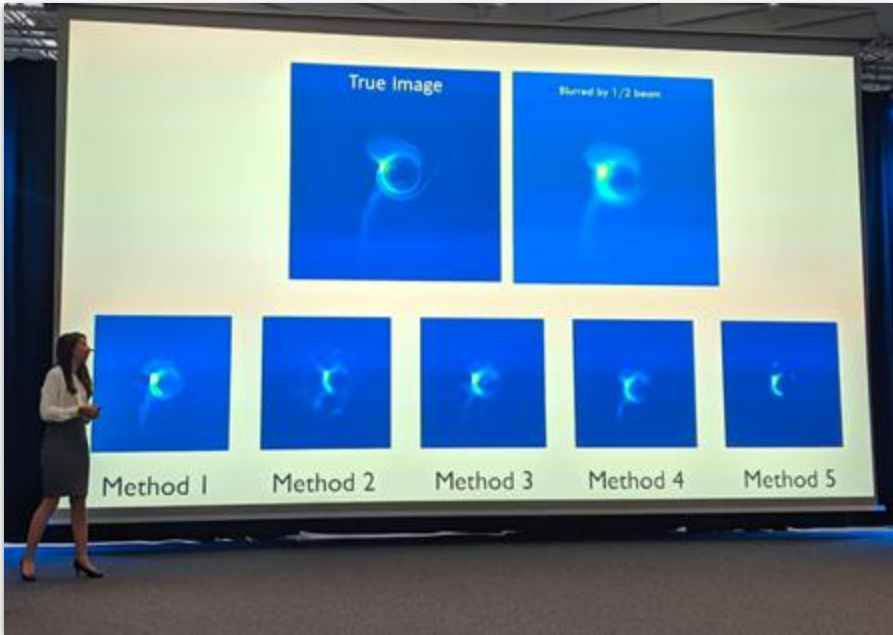
The Event Horizon Telescope Collaboration, Kazunori Akiyama<sup>1,2,3,4</sup>, Anthon Alberdi<sup>5</sup>, Walter Alef<sup>6</sup>, Keiichi Asada<sup>7</sup>, Rebecca Azulay<sup>8,9,6</sup>, Anne-Kathrin Baczko<sup>6</sup>, David Ball<sup>10</sup>, Mislav Balokovic<sup>4,11</sup>, John Barrett<sup>2</sup>, [Show full author list](#)  
Published 2019 April 10 • © 2019, The American Astronomical Society.  
[The Astrophysical Journal Letters, Volume 875, Number 1](#)

Software: DiFX (Deller et al. 2011), CALC, PoIConvert (Marti-Vidal et al. 2016), HOPS (Whitney et al. 2004), CASA (McMullin et al. 2007), AIPS (Greisen 2003), ParselTongue (Kettenis et al. 2006), GNU Parallel (Tange 2011), GILDAS, eht-imaging (Chael et al. 2016, 2018), Numpy (van der Walt et al. 2011), Scipy (Jones et al. 2001), Pandas (McKinney 2010), Astropy (The Astropy Collaboration et al. 2013, 2018), Jupyter (Kluyver et al. 2016), Matplotlib (Hunter 2007).

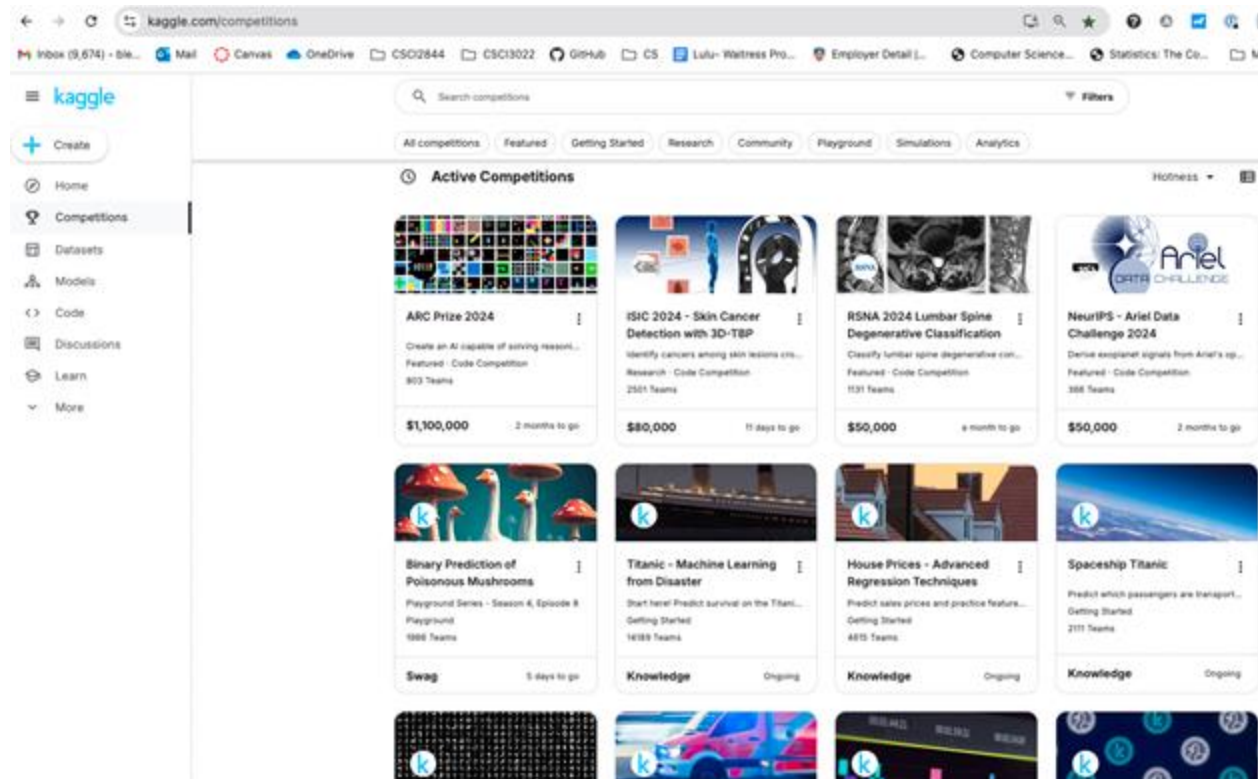
22,953  
Community Contributors

Analysis Courtesy of GitHub

# 2019: First Image of a Black Hole



# Kaggle: More Data Science Sets and Competitions



<https://www.kaggle.com/competitions>



# List of Topics to be Covered in CSCI 3022

- **Exploring, Cleaning and Visualizing Data**

- Intro to Pandas and NumPy
- Exploratory Data Analysis
- Wrangling Data
- Visualizing Data using matplotlib, seaborn & plotly

- **Probability & Statistics for Data Science**

- Independence
- Conditioning and Bayes Theorem
- Discrete and Continuous Random Variables
- Distributions and Joint Distributions
- Expectation & Variance
- Central Limit Theorem
- Sampling
- Using Statistical Simulation To Draw Inferences from Data:
  - Hypothesis and A/B Testing
  - Bootstrapping Confidence Intervals

- **Modeling/Intro to Machine Learning**

- Model design and loss formulation
- Simple Linear Regression
- Multiple Linear Regression
- Logistic Regression
- Feature Engineering
- Cross-Validation
- Regularization

**iclicker Poll:** <https://join.iclicker.com/XSFZ>

**Which course topic are you most interested to learn?**



## Lesson Learning Objectives:

- Name and explain the stages of the data science lifecycle

# The Data Science Lifecycle

---

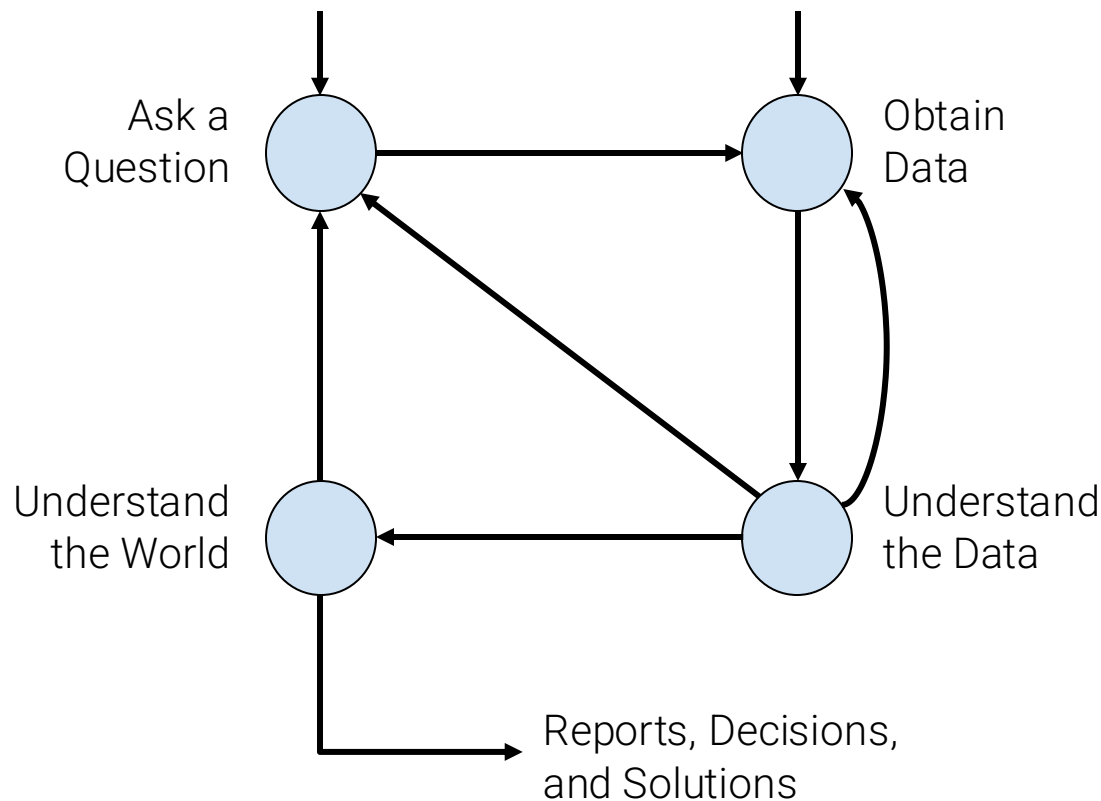
Lecture 01, CSCI 3022

- Intros & Logistics
- What is Data Science?
  - **Data Science Lifecycle**
- Exploratory Data Analysis & Wrangling
- Intro to Pandas
  - Jupyter Notebook Demo: EDA using Pandas

# Data science lifecycle

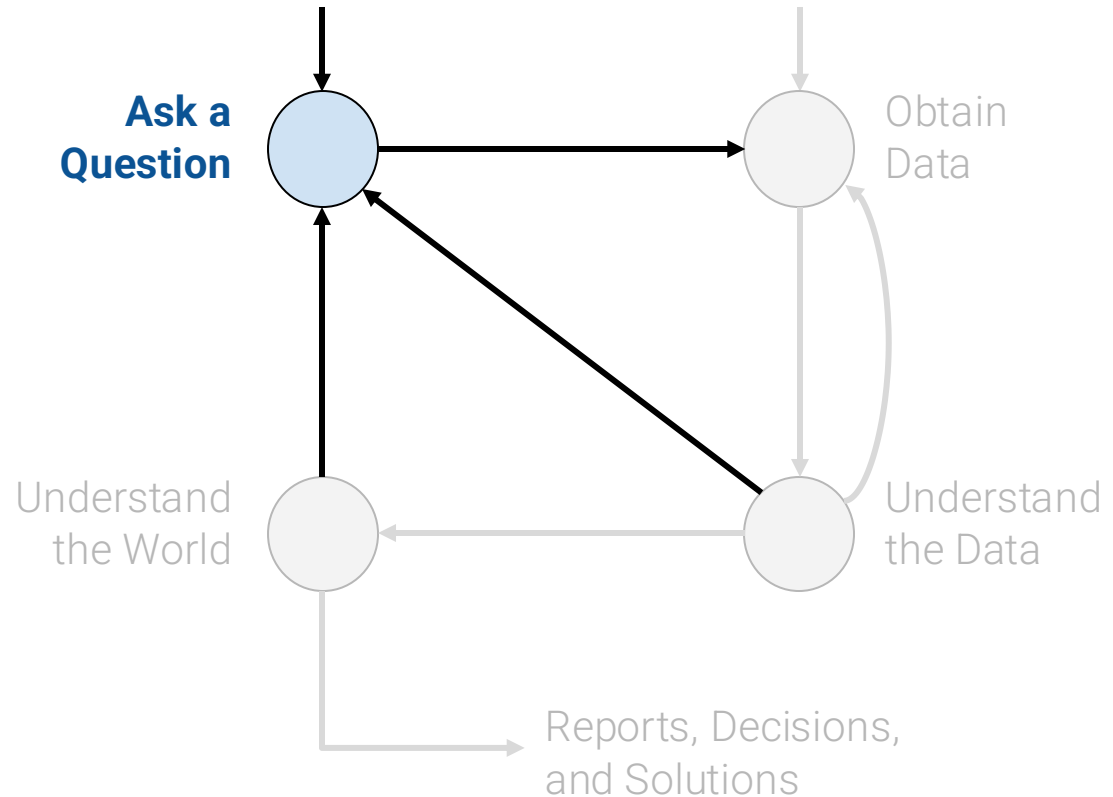
The data science lifecycle is a **high-level description** of the data science workflow.

Note the two distinct entry points!



# 1. Question/Problem Formulation

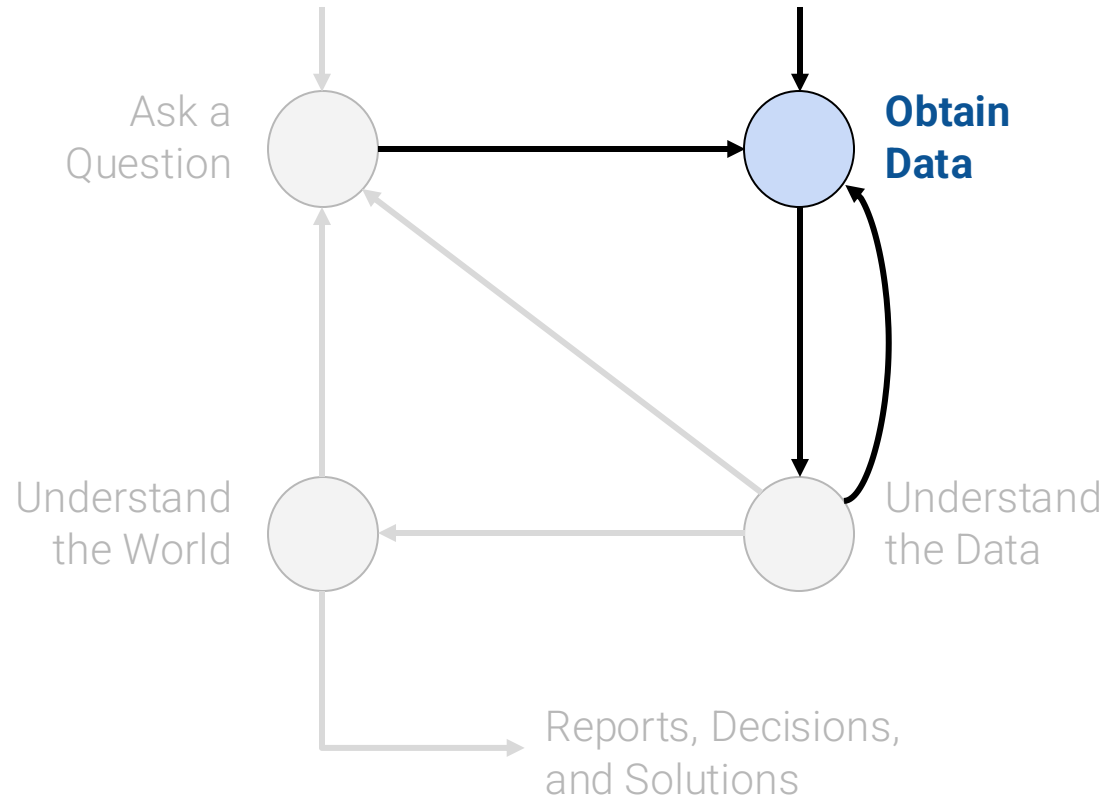
- What do we want to know?
- What problems are we trying to solve?
- What are the hypotheses we want to test?
- What are our metrics for success?





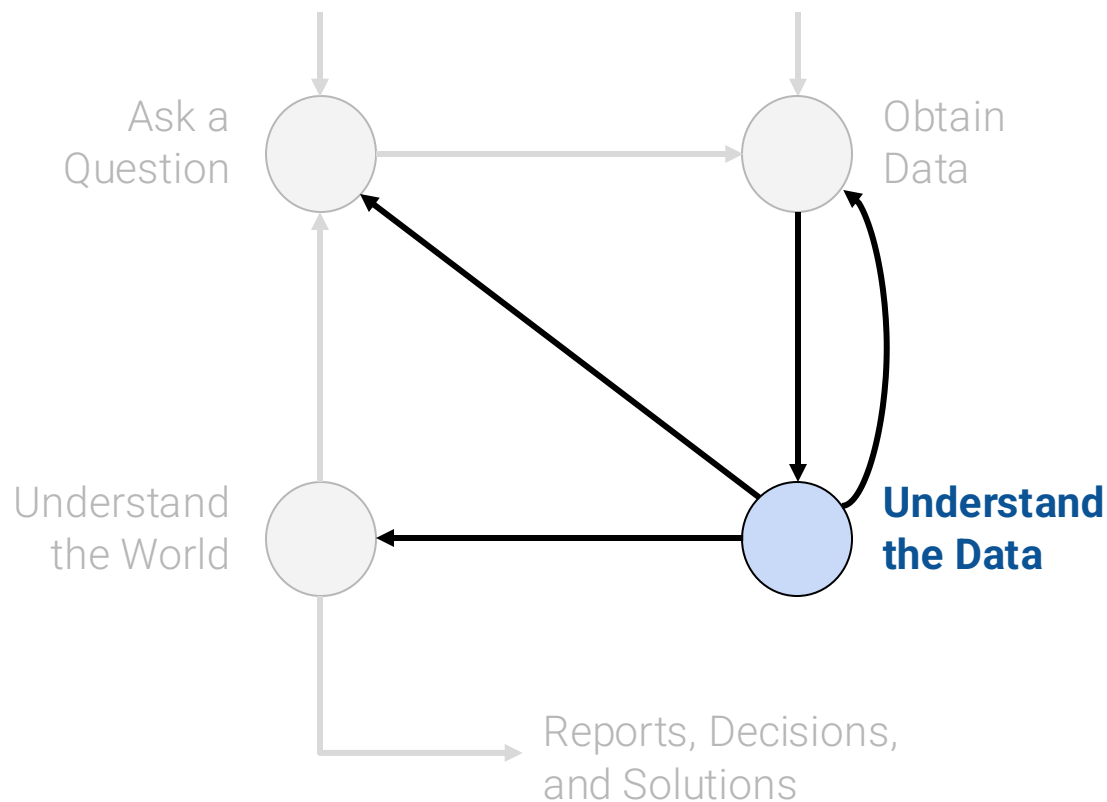
## 2. Data Acquisition and Cleaning

- What data do we have and what data do we need?
- How will we sample more data?
- Is our data representative of the population we want to study?



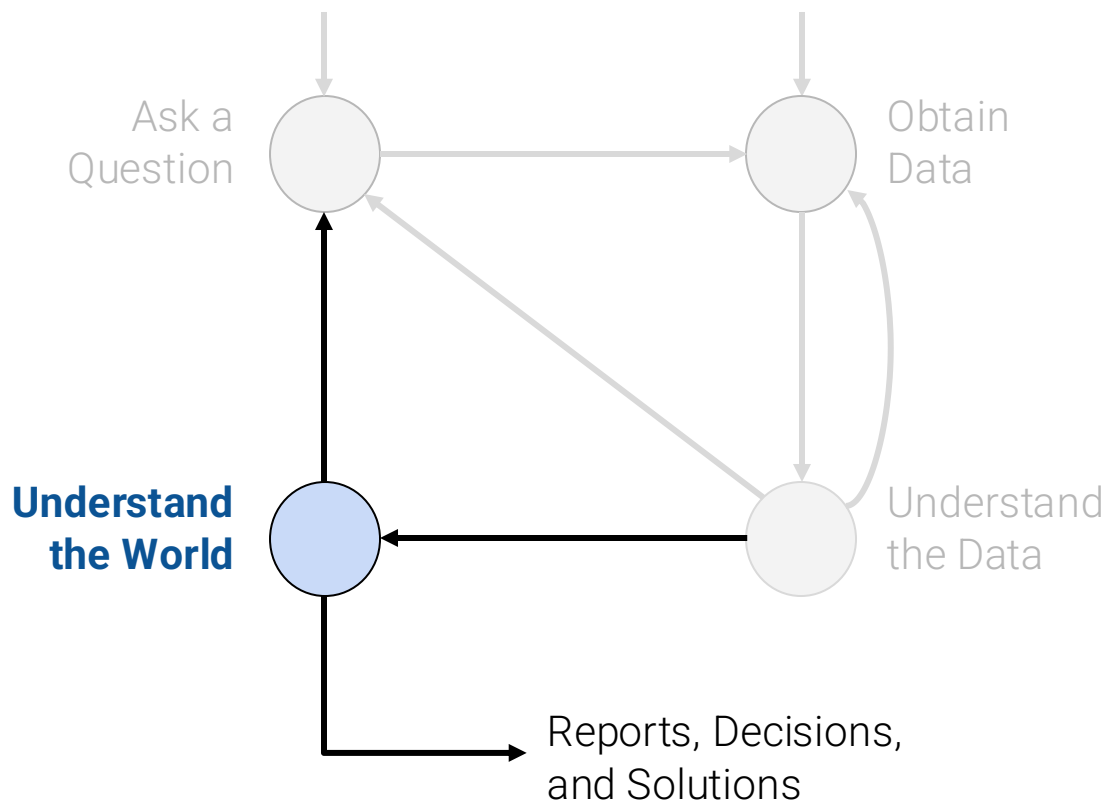
### 3. Exploratory Data Analysis & Visualization

- How is our data organized and what does it contain?
- Do we already have relevant data?
- What are the biases, anomalies, or other issues with the data?
- How do we transform the data to enable effective analysis?




## 4. Prediction and Inference

- What does the data say about the world?
- Does it answer our questions or accurately solve the problem?
- How robust are our conclusions and can we trust the predictions?



# Plan for first few weeks



**(Weeks 1 and 2)**  
EDA, Wrangling and Data Visualization

## Lesson Learning Objectives:

- Meet your classmates
- Name and explain the stages of the data science lifecycle
- Identify 5 key data properties to consider when doing Exploratory Data Analysis and implement using demo data

# EDA & Wrangling

---

Lecture 01, CSCI 3022

- Intros & Logistics
- What is Data Science?
- **Exploratory Data Analysis & Wrangling**
- Intro to Pandas:
  - Jupyter Notebook Demo: EDA using Pandas



## Congratulations!!!

You **have collected** or **have been given** a box of data.

What do you do next?

## One Option: Exploratory Data Analysis (EDA)

*"Getting to know and understand the data"*

The process of **transforming**, **visualizing**, and **summarizing** data to:

- Build/confirm understanding of the data and its **provenance**
- Identify and address potential issues in the data.
- Inform the subsequent analysis.
- Discover *potential* hypotheses ...

*Provenance*: origin of data;  
methodology by which data  
were produced

**EDA is an open-ended analysis.**

- Informal, no specific idea of what we are looking for.
- Be willing to find something surprising!

Contrast with **confirmatory analysis**:

- Questions are fixed in advance.
- Allows for more rigorous statistical analysis.



## Another Option: Data Wrangling (Data Cleaning)

### Data Wrangling, or Data Cleaning:

The process of transforming **raw data** to facilitate subsequent analysis.


Often addresses **issues** like...

- structure / formatting
- missing or corrupted values
- unit conversion
- encoding text as numbers
- ...

Sadly, data cleaning is a big part of data science...









**EDA is unboxing  
for data!**

## Exploratory Data Analysis (EDA) Guiding Principles

*“Exploratory data analysis is an attitude, a state of flexibility, a willingness to look for those things that we believe are not there, as well as those that we believe to be there.”* – John Tukey

# Key Data Properties to Consider in EDA

**Structure** -- the “shape” of a data file


**Granularity** -- what does each record represent?

**Scope** -- how (in)complete is the data

**Temporality** -- how is the data situated in time

**Faithfulness** -- how well does the data capture “reality”

File Format  
Variable Type  
Multiple files  
(Primary and Foreign Keys)



**Structure** -- the “shape” of a data file

**Granularity** -- how fine/coarse is each datum

**Scope** -- how (in)complete is the data

**Temporality** -- how is the data situated in time


**Faithfulness** -- how well does the data capture “reality”

# Rectangular and Non-rectangular Data



Data come in many different shapes.

Dataset's *structure*: Mental representation of the data

Tabular (rectangular) data



Non-rectangular data




# Tabular/Rectangular Data

We often prefer tabular data for data analysis (why?)

- Regular structures are easy to manipulate and analyze
- A big part of data cleaning is about transforming data to be more rectangular

Two kinds of tabular data: **Tables** and **Matrices**.



**Tables** (a.k.a. **DataFrames** in R/Python and relations in SQL)

- Named columns with different types
- Manipulated using Pandas data transformation languages (map, filter, group by, join, ...)

**Matrices**

- Numeric data of the same type (float, int, etc.)
- Manipulated using linear algebra

# Data Scientists Love Rectangular/Tabular Data

"Tabular= data in a table.

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
177	2016	Jill Stein	Green	1457226	loss	1.073699
178	2020	Joseph Biden	Democratic	81268924	win	51.311515
179	2020	Donald Trump	Republican	74216154	loss	46.858542
180	2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
181	2020	Howard Hawkins	Green	405035	loss	0.255731

A **row** represents one record (i.e. an observation)

A **column** represents some characteristic, or feature, of that observation (here, the political party of that person).

A tabular **dataset** is **tidy** when each column corresponds to one **variable** in the **dataset**, each row corresponds to one **observation**, and all **variables** in the **dataset** have the same **unit of observation**.

# Intro To Pandas

---

Lecture 01, CSCI 3022

- Intros & Logistics
- What is Data Science?
- Exploratory Data Analysis & Wrangling
- Intro To Pandas
  - Jupyter Notebook Demo: EDA using Pandas



# Introducing the Standard Python Data Science Tool: Pandas

---

The Python Data  
Analysis Library



```
# data manipulation  
import pandas as pd
```

- **Pandas** (derived from Panel Data) is a Data Analysis library to make data cleaning and analysis fast and convenient in Python.
- Pandas adopts many coding idioms from NumPy, the biggest difference is that ***pandas is designed for working with tabular or heterogeneous data.***
  - NumPy by contrast is best suited for working with homogenous numerical data.

Tabular data is one of the most common data formats.

- Will be our primary focus in CSCI 3022

## Introducing the Standard Python Data Science Tool: pandas

---

Using `pandas`, we can:

- Arrange data in a tidy tabular format.
- Extract useful information filtered by specific conditions.
- Operate on data to gain new insights.
- Apply NumPy functions to our data
- Perform vectorized computations to speed up our analysis.

`pandas` is the standard tool across research and industry for working with tabular data.

The first week of this course will serve as a "bootcamp" in helping you build familiarity with operating on data with `pandas`.

# Pandas Data Structures

There are three fundamental data structures in pandas:

- **Series**: 1D labeled array data. I usually think of it as columnar data.
- **Data Frame**: 2D tabular data with both row and column labels
- **Index**: A sequence of row/column labels.

Data Frame

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

Series named "Candidate"

0	Obama
1	McCain
2	Obama
3	Romney
4	Clinton
5	Trump
Name: Candidate, dtype: object	

Index

# The Relationship Between DataFrames, Series, and Indices

We can think of a **DataFrame** as a collection of **Series** that all share the same **Index**.

- Candidate, Party, %, Year, and Result **Series** all share an **Index** from 0 to 5.

Candidate Series    Party Series    % Series    Year Series    Result Series



	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

## Lesson Learning Objectives:

- Identify 5 key data properties to consider when doing Exploratory Data Analysis and implement using demo data

# Jupyter Demo

---

Lecture 01, CSCI 3022

- Intros & Logistics
- What is Data Science?
- Exploratory Data Analysis & Wrangling
- Intro to Pandas
  - Jupyter Notebook Demo: EDA using Pandas

## CSV: Comma-Separated Values

Election Data in the US

CSV is a very common **tabular file format**.

- **Records** (rows) are delimited by a newline: `'\n'`, `"\r\n"`
- **Fields** (columns) are delimited by commas: `' , '`

Pandas: `pd.read_csv(header=...)`

## Demo Slides

Records/Rows	Fields/Attributes/Features/Columns		
		Year	Candidate
	0	2024	Kamala Harris
	1	2024	Donald Trump

# Creating a DataFrame

Many approaches exist for creating a DataFrame.

## 1). From a CSV file.

```
elections = pd.read_csv("data/elections.csv")
```

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
177	2016	Jill Stein	Green	1457226	loss	1.073699
178	2020	Joseph Biden	Democratic	81268924	win	51.311515
179	2020	Donald Trump	Republican	74216154	loss	46.858542
180	2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
181	2020	Howard Hawkins	Green	405035	loss	0.255731

The DataFrame elections

```
elections = pd.read_csv("data/elections.csv", index_col="Year")
```

	Candidate	Party	Popular vote	Result	%
Year					
1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
1828	Andrew Jackson	Democratic	642806	win	56.203927
1828	John Quincy Adams	National Republican	500897	loss	43.796073
1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...
2016	Jill Stein	Green	1457226	loss	1.073699
2020	Joseph Biden	Democratic	81268924	win	51.311515
2020	Donald Trump	Republican	74216154	loss	46.858542
2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
2020	Howard Hawkins	Green	405035	loss	0.255731

The DataFrame elections with "Year" as Index



# Supporting Materials

---

Lecture 01, CSCI 3022

**Supporting Materials: Pandas Data Structures:**

**Series**

**DataFrames**

**Indices**



## Learning Objectives

- Understand the relationship between Series, DataFrames and Indices in Pandas
- Create, filter and perform operations on Series

# Pandas Data Structures: Series

---

- Introduction to Series

# Pandas Data Structures

There are three fundamental data structures in pandas:

- **Series**: 1D labeled array data. I usually think of it as columnar data.
- **Data Frame**: 2D tabular data with both row and column labels
- **Index**: A sequence of row/column labels.

Data Frame

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

Series named "Candidate"

0	Obama
1	McCain
2	Obama
3	Romney
4	Clinton
5	Trump
Name: Candidate, dtype: object	

Index

# Series

A **Series** is a 1-dimensional array-like object. It contains:

- A sequence of **values** of the same type.
- A sequence of data labels, called the **index**.

**pd** is the conventional alias for **pandas**

```
import pandas as pd  
s = pd.Series(["welcome", "to", "cspb 3022"])
```

0	welcome
1	to
2	CSPB 3022

dtype: object

**Index**, accessed by calling `s.index`

**Values**, accessed by calling `s.values`

## Series - Custom Index

- We can provide index labels for items in a Series by passing an index list.

```
s = pd.Series([-1, 10, 2], index = ["a", "b", "c"])
```

```
a    -1  
b    10  
c     2  
dtype: int64
```

```
s.index
```

```
Index(['a', 'b', 'c'], dtype='object')
```

- A Series index can also be changed.

```
s.index = ["first", "second", "third"]
```

```
first    -1  
second   10  
third     2  
dtype: int64
```

```
s.index
```

```
Index(['first', 'second', 'third'], dtype='object')
```

## Selection in Series

- We can select a single value or a set of values in a Series using:
  - A single label
  - A list of labels
  - A filtering condition

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a      4
b     -2
c      0
d      6
dtype: int64
```

## Selection in Series

- We can select a single value or a set of values in a Series using:
  - **A single label**
  - A list of labels
  - A filtering condition

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a    4
b   -2
c    0
d    6
dtype: int64
```

```
s["a"]
```

```
4
```

## Selection in Series

- We can select a single value or a set of values in a Series using:
  - A single label
  - **A list of labels**
  - A filtering condition

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a    4
b   -2
c    0
d    6
dtype: int64
```

```
s[["a", "c"]]
```

```
a    4
c    0
dtype: int64
```

## Selection in Series

- We can select a single value or a set of values in a **Series** using:
  - A single label
  - A list of labels
  - **A filtering condition**

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a    4
b   -2
c    0
d    6
dtype: int64
```

- Say we want to select values in the **Series** that satisfy a particular condition:
  - 1) Apply a boolean condition to the **Series**. This creates a **new Series of boolean values**.
  - 2) Index into our **Series** using this boolean condition. **pandas** will select only the entries in the **Series** that satisfy the condition.

```
s > 0
```

```
a    True
b   False
c   False
d    True
dtype: bool
```

```
s[s > 0]
```

```
a    4
d    6
dtype: int64
```



## Practice Question

---

What is the output of the following code?

```
example = pd.Series([4, 5, 6], index=["one", "two", "three"])
example[example > 4].values
```

A). `array([4, 5, 6])`

C). `array([5, 6])`

`Index(['two', 'three'], dtype='object')`

D). `[5, 6]`

B).

E). None of these

## Practice Question

---

What is the output of the following code?

```
example = pd.Series([4, 5, 6], index=["one", "two", "three"])  
example[example > 4].values
```

A). `array([4, 5, 6])`

`Index(['two', 'three'], dtype='object')`

B).

C). `array([5, 6])`

D). `[5, 6]`

E). None of these

# DataFrames of Series!

Typically, we will work with **Series** using the perspective that they are columns in a **DataFrame**.

We can think of a **DataFrame** as a collection of **Series** that all share the same **Index**.

0

1824

1

1824

2

1828

3

1828

4

1832

...

177

2016

178

2020

179

2020

180

2020

181

2020

Name: Year,

0

Andrew Jackson

1

John Quincy Adams

2

Andrew Jackson

3

John Quincy Adams

4

Andrew Jackson

...

177

Jill Stein

178

Joseph Biden

179

Donald Trump

180

Jo Jorgensen

181

Howard Hawkins

Name: Candidate,

+

[...]

→

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
177	2016	Jill Stein	Green	1457226	loss	1.073699
178	2020	Joseph Biden	Democratic	81268924	win	51.311515
179	2020	Donald Trump	Republican	74216154	loss	46.858542
180	2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
181	2020	Howard Hawkins	Green	405035	loss	0.255731

The Series "Year"

The Series "Candidate"

The DataFrame elections



Non-native English speaker note: The plural of "series" is "series". Sorry.

## Learning Objectives

- Understand the relationship between DataFrames and Indices in Pandas
- Create DataFrames
- Manipulate indices

# Pandas Data Structures

---

- Introduction to DataFrames
- Indices

# The Relationship Between DataFrames, Series, and Indices

We can think of a **DataFrame** as a collection of **Series** that all share the same **Index**.

- Candidate, Party, %, Year, and Result **Series** all share an **Index** from 0 to 5.

Candidate Series    Party Series    % Series    Year Series    Result Series



	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

# The DataFrame API

---

The API for the **DataFrame** class is enormous.

- API: "Application Programming Interface".
- The API is the set of abstractions supported by the class.

Full documentation is at

<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>

- We will only consider a tiny portion of this API.

We want you to get familiar with the real world programming practice of... Googling!

- Answers to your questions are often found in the **pandas** documentation, Stack Overflow, etc.

## Creating a DataFrame

---

The syntax of creating **DataFrame** is:

```
pandas.DataFrame(data, index, columns)
```

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- Using a list and column name(s).
- From a dictionary.
- From a **Series**.

# Creating a DataFrame

The syntax of creating `DataFrame` is:

```
pandas.DataFrame(data, index, columns)
```

Many approaches exist for creating a `DataFrame`. Here, we will go over the most popular ones.

- **From a CSV file.**
- Using a list and column name(s).
- From a dictionary.
- From a `Series`.

```
elections = pd.read_csv("data/elections.csv")
```

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
177	2016	Jill Stein	Green	1457226	loss	1.073699
178	2020	Joseph Biden	Democratic	81268924	win	51.311515
179	2020	Donald Trump	Republican	74216154	loss	46.858542
180	2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
181	2020	Howard Hawkins	Green	405035	loss	0.255731

The `DataFrame` `elections`



# Creating a DataFrame

The syntax of creating `DataFrame` is:

```
pandas.DataFrame(data, index, columns)
```

Many approaches exist for creating a `DataFrame`. Here, we will go over the most popular ones.

- **From a CSV file.** `elections = pd.read_csv("data/elections.csv", index_col="Year")`
- Using a list and column name(s).
- From a dictionary.
- From a `Series`.

	Candidate	Party	Popular vote	Result	%
Year					
1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
1828	Andrew Jackson	Democratic	642806	win	56.203927
1828	John Quincy Adams	National Republican	500897	loss	43.796073
1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...
2016	Jill Stein	Green	1457226	loss	1.073699
2020	Joseph Biden	Democratic	81268924	win	51.311515
2020	Donald Trump	Republican	74216154	loss	46.858542
2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
2020	Howard Hawkins	Green	405035	loss	0.255731

The `DataFrame` `elections` with `"Year"` as Index

## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- **Using a list and column name(s).**
- From a dictionary.
- From a **Series**.

```
pd.DataFrame([1, 2, 3],  
             columns=["Numbers"])
```

Numbers	
0	1
1	2
2	3

```
pd.DataFrame([[1, "one"], [2, "two"]],  
             columns = ["Number", "Description"])
```

Number Description		
0	1	one
1	2	two

## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- Using a list and column name(s).
- **From a dictionary.**
- From a **Series**.

Wait, what's a dictionary?

**Dictionary in Python**  PYnative.com  
Unordered collections of unique values stored in (Key-Value) pairs.

```
d = {'a': 10, 'b': 20, 'c': 30}
```

↑      ↑      ↑  
d['a'] d['b'] d['c']

- ✓ **Unordered:** The items in dict are stored without any index value
- ✓ **Unique:** Keys in dictionaries should be Unique
- ✓ **Mutable:** We can add/Modify/Remove key-value after the creation


```
d2={"Fruit":["Strawberry",  
"Orange"], "Price": [5.49, 3.99]}
```

## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.


- From a CSV file.
- Using a list and column name(s).
- **From a dictionary.**
- From a **Series**.

```
pd.DataFrame({"Fruit":["Strawberry", "Orange"],  
             "Price": [5.49, 3.99]})
```



Specify columns of the **DataFrame**

```
pd.DataFrame([{"Fruit":"Strawberry", "Price":5.49},  
             {"Fruit":"Orange", "Price":3.99}])
```



Specify rows of the **DataFrame**

	Fruit	Price
0	Strawberry	5.49
1	Orange	3.99

## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- Using a list and column name(s).
- From a dictionary.
- **From a Series.**

```
s_a = pd.Series(["a1", "a2", "a3"], index = ["r1", "r2", "r3"])  
s_b = pd.Series(["b1", "b2", "b3"], index = ["r1", "r2", "r3"])
```

```
pd.DataFrame({"A-column":s_a, "B-column":s_b})
```

	A-column	B-column
r1	a1	b1
r2	a2	b2
r3	a3	b3

```
pd.DataFrame(s_a)
```

```
s_a.to_frame()
```

	0
r1	a1
r2	a2
r3	a3

# Indices Are Not Necessarily Row Numbers

A **Row Index** (a.k.a. row labels) can also:

- Be non-numeric.
- Have a name, e.g. "Candidate".

```
# Creating a DataFrame from a CSV file and specifying the Index column
elections = pd.read_csv("data/elections.csv", index_col = "Candidate")
```

Candidate	Year	Party	Popular vote	Result	%
Andrew Jackson	1824	Democratic-Republican	151271	loss	57.210122
John Quincy Adams	1824	Democratic-Republican	113142	win	42.789878
Andrew Jackson	1828	Democratic	642806	win	56.203927
John Quincy Adams	1828	National Republican	500897	loss	43.796073
Andrew Jackson	1832	Democratic	702735	win	54.574789

# Indices Are Not Necessarily Unique

The row labels that constitute an index do not have to be unique.

- Left: The **index** values are all unique and numeric, acting as a row number.
- Right: The **index** values are named and non-unique.

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

Year	Candidate	Party	%	Result
2008	Obama	Democratic	52.9	win
2008	McCain	Republican	45.7	loss
2012	Obama	Democratic	51.1	win
2012	Romney	Republican	47.2	loss
2016	Clinton	Democratic	48.2	loss
2016	Trump	Republican	46.1	win

# Modifying Indices

- We can select a new column and set it as the index of the **DataFrame**.

Example: Setting the index to the "Party" column.

```
elections.set_index("Party")
```

	Candidate	Year	Popular vote	Result	%
Party					
Democratic-Republican	Andrew Jackson	1824	151271	loss	57.210122
Democratic-Republican	John Quincy Adams	1824	113142	win	42.789878
Democratic	Andrew Jackson	1828	642806	win	56.203927
National Republican	John Quincy Adams	1828	500897	loss	43.796073
Democratic	Andrew Jackson	1832	702735	win	54.574789
...	...	...	...	...	...
Green	Jill Stein	2016	1457226	loss	1.073699
Democratic	Joseph Biden	2020	81268924	win	51.311515
Republican	Donald Trump	2020	74216154	loss	46.858542
Libertarian	Jo Jorgensen	2020	1865724	loss	1.177979
Green	Howard Hawkins	2020	405035	loss	0.255731






# Resetting the Index

- We can change our mind and reset the **Index** back to the default list of integers.

```
elections.reset_index()
```

	Candidate	Year	Popular vote	Result	%
Party					
Democratic-Republican	Andrew Jackson	1824	151271	loss	57.210122
Democratic-Republican	John Quincy Adams	1824	113142	win	42.789878
Democratic	Andrew Jackson	1828	642806	win	56.203927
National Republican	John Quincy Adams	1828	500897	loss	43.796073
Democratic	Andrew Jackson	1832	702735	win	54.574789
...	...	...	...	...	...
Green	Jill Stein	2016	1457226	loss	1.073699
Democratic	Joseph Biden	2020	81268924	win	51.311515
Republican	Donald Trump	2020	74216154	loss	46.858542
Libertarian	Jo Jorgensen	2020	1865724	loss	1.177979
Green	Howard Hawkins	2020	405035	loss	0.255731

	Candidate	Year	Party	Popular vote	Result	%
0	Andrew Jackson	1824	Democratic-Republican	151271	loss	57.210122
1	John Quincy Adams	1824	Democratic-Republican	113142	win	42.789878
2	Andrew Jackson	1828	Democratic	642806	win	56.203927
3	John Quincy Adams	1828	National Republican	500897	loss	43.796073
4	Andrew Jackson	1832	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
177	Jill Stein	2016	Green	1457226	loss	1.073699
178	Joseph Biden	2020	Democratic	81268924	win	51.311515
179	Donald Trump	2020	Republican	74216154	loss	46.858542
180	Jo Jorgensen	2020	Libertarian	1865724	loss	1.177979
181	Howard Hawkins	2020	Green	405035	loss	0.255731



## Column Names Are Usually Unique!

---

Column names in **pandas** are almost always unique.

- Example: Really shouldn't have two columns named "Candidate".

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

# Accessing a DataFrame's columns and row indices: .index and .columns

```
elections = pd.read_csv("data/elections.csv", index_col="Year")
```

Year	Candidate	Party	Popular vote	Result	%
1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
1828	Andrew Jackson	Democratic	642806	win	56.203927
1828	John Quincy Adams	National Republican	500897	loss	43.796073
1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...
2016	Jill Stein	Green	1457226	loss	1.073699
2020	Joseph Biden	Democratic	81268924	win	51.311515
2020	Donald Trump	Republican	74216154	loss	46.858542
2020	Jo Jorgensen	Libertarian	1865724	loss	1.177979
2020	Howard Hawkins	Green	405035	loss	0.255731

```
elections.index
```

```
Index([1824, 1824, 1828, 1828, 1832, 1832, 1832, 1836, 1836, 1836,
      ...,
      2016, 2016, 2016, 2016, 2016, 2016, 2020, 2020, 2020, 2020],
      dtype='int64', name='Year', length=182)
```

```
elections.columns
```

```
Index(['Candidate', 'Party', 'Popular vote', 'Result', '%'], dtype='object')
```

The DataFrame `elections` with `"Year"` as Index

